

# vailability That Instantly Recruiters

BY DEVOPS SHACK



### **DevOps Shack**

## **High Availability Skills That Instantly Impress Recruiters**

- Basic "HA Web Server with NGINX & Keepalived on AWS EC2"
  - Build a basic two-node web server cluster with automatic failover using Keepalived for VIP management.
- - Deploy Jenkins in HA mode with multiple replicas, shared NFS volume for /var/jenkins\_home, and Kubernetes Horizontal Pod Autoscaling.
- Advanced "Production-Grade HA Microservices Stack with Istio, Vault, MongoDB, and ArgoCD on EKS"
  - Full-scale project deploying multi-tier apps with HA MongoDB StatefulSet, HA Vault with Raft, Canary rollout using ArgoCD, and Istio- managed routing.



#### High Availability in DevOps Projects

# What Recruiters Love to Hear — From Basics to Production-Grade Architectures

In today's digital-first world, **downtime is expensive** — every second of unavailability can cost businesses users, revenue, and trust. That's why **High Availability (HA)** is no longer optional. It's a **core skill** every serious DevOps engineer must master.

But here's the twist: while many talk about HA, **few know how to design and implement it practically**. This guide fills that gap with **real-world DevOps projects** that showcase your ability to build resilient, fault-tolerant, self-healing systems.

Whether you're a beginner just getting started, an intermediate learner scaling your skills, or an advanced engineer prepping for top-tier interviews — this guide walks you through **3 real HA projects**, each one designed to:

=	Build confidence in designing, deploying, and explaining fault-tolerant systems
	<b>Why It Matters:</b>
	High Availability (HA) isn't just a buzzword — it's <b>business-critical</b> . Recruiters <b>love candidates</b> who:
	<ul> <li>✓ Think in terms of fault tolerance, not just features</li> <li>✓ Architect for 99.99% uptime or better</li> <li>✓ Proactively design recovery, not just deployment</li> </ul>
٠,	hether you're applying for a DevOps SPE or Platform Engineer role — th

Deepen your technical architecture knowledge

Whether you're applying for a DevOps, SRE, or Platform Engineer role — the ability to think and build for resilience is a skill that sets you apart.

Let's break it down into what you can **showcase in interviews, resumes, or projects** to highlight HA experience:

#### What Is High Availability (HA)?

High Availability means your systems remain **operational and accessible** even when:



#### A pod or container crashes

- A node goes down
- · A region becomes unavailable
- A disk fails or IOPS drops
- A service gets overwhelmed with

traffic In technical terms, HA requires:

- Redundancy: Multiple instances of critical components
- Load Balancing: Traffic distribution to healthy nodes
- Failover Mechanisms: Automated switching to backup systems
- Health Checks: Probes to detect and react to failures
- State Management: Persistent data storage that survives failure

Recruiters don't just want to know that you can *run* apps—they want to see that you can keep them running **no matter what**.

**☑** Project 1: High Availability Web Server with NGINX & Keepalived on AWS EC2





Level:

Basic Street Objective:

Deploy two NGINX web servers behind a **Virtual IP (VIP)** using **Keepalived** to ensure **automatic failover** when one EC2 instance goes down.

#### What You'll Learn

- Basic High Availability concepts using active-passive setup
- How to install and configure NGINX
- How to configure Keepalived to manage a floating IP
- Testing automatic failover when a node goes down

#### **Architecture Overview**

**Note:** Only one instance holds the VIP at a time. If the active one fails, the passive node becomes active.

#### Prerequisites

 2 AWS EC2 instances (Ubuntu 22.04 recommended) in the same VPC and subnet



#### A secondary private IP address in the subnet to act as a floating VIP

- SSH access with sudo privileges to both instances
- Security Group allowing:
  - TCP port 80 (HTTP)
  - Protocol 112 (VRRP)
  - ICMP (for health check pings)
- Step 1: Allocate a Secondary Private IP (VIP)
  - 1. Go to EC2 Dashboard → Network Interfaces
  - 2. Find the network interface attached to Instance A
  - 3. Click Actions → Manage IP Addresses → Assign new IP
  - 4. Note down the new private IP (e.g., 192.168.1.100)
  - 5. **Do NOT associate this with any instance manually** Keepalived will do this dynamically.
- Step 2: Install NGINX & Keepalived

Run the following on **both instances**:

sudo apt update

sudo apt install -y nginx keepalived net-tools

Check installation:

nginx -v keepalived -

٧

Step 3: Configure NGINX

On **both nodes**, replace the default NGINX index page:

echo "Welcome to Instance A" | sudo tee /var/www/html/index.html #

Or on B:





echo "Welcome to Instance B" | sudo tee /var/www/html/index.html

```
Restart NGINX:
```

```
sudo systemctl restart nginx

Check on browser: http://<EC2_IP>
```

#### Step 4: Configure Keepalived on Instance A (MASTER)

```
Create the config file:
```

```
sudo nano /etc/keepalived/keepalived.conf
```

#### Paste:

```
vrrp_instance VI_1
    { state MASTER
    interface eth0
    virtual_router_id 51
    priority 101
    advert_int 1
    authentication {
        auth_type PASS
        auth_pass devops123
    }
    virtual_ipaddress {
        192.168.1.100
    }
}
Save and restart:
sudo systemctl restart
```

keepalived Check:



```
ip a | grep 192.168.1.100
```

**✓** Check active instance:

You should see the VIP bound to eth0 on Instance A.

```
Step 5: Configure Keepalived on Instance B (BACKUP)
Edit config:
sudo nano /etc/keepalived/keepalived.conf
Paste:
vrrp_instance VI_1
  { state BACKUP
  interface eth0
  virtual_router_id 51
  priority 100
  advert_int 1
  authentication {
    auth_type PASS
    auth_pass devops123
  }
  virtual_ipaddress {
    192.168.1.100
  }
Restart service:
sudo systemctl restart keepalived
Instance B won't show the VIP yet (expected).
   Step 6: Test High Availability (Failover)
```



#### curl

# Should show: Welcome to Instance A

Simulate failure:

On Instance A:

sudo systemctl stop keepalived

On Instance B:

ip a | grep 192.168.1.100

# Should now see VIP assigned to eth0

Test again:

curl http://192.168.1.100

# Output: Welcome to Instance B

Success! HA failover works.

- Security Consideration
  - Ensure VRRP (protocol 112) is allowed in the security group
  - Use stronger auth pass for production
  - Disable root SSH if not required
  - Consider health check scripts and notifications for production
- What to Mention in Resume or Interview
- ✓ Built an HA web server cluster using NGINX & Keepalived
- Configured VIP failover between EC2 instances via VRRP protocol
- Designed Active-Passive topology to simulate real-world disaster recovery
- ✓ Validated system resilience via simulated node failure tests
- **ℒ** Project 2: Highly Available Jenkins CI/CD with NFS Shared Storage on Kubernetes

Level:

Intermediate 🖄

**Objective:** 





Deploy **Jenkins in HA mode** on Kubernetes using **multiple replicas** backed by a **shared NFS volume**, ensuring that CI/CD pipelines stay available even if a pod crashes or gets rescheduled.

#### **What You'll Learn**

- Deploy Jenkins in **High Availability mode** on Kubernetes
- Configure Persistent Shared Storage using NFS
- Use ReadWriteMany (RWX) volumes to support multiple replicas
- Leverage Horizontal Pod Autoscaler (HPA) and liveness probes
- Test pod crash recovery and job state persistence

#### **Architecture Overview**

#### Prerequisites

- Kubernetes cluster (minikube or EKS/GKE/AKS)
- NFS Server deployed (or use NFS provisioner)
- Helm installed





#### kubectl configured

Ingress controller (NGINX or ALB Ingress)

#### Step 1: Setup a Shared NFS Volume (Using NFS Provisioner)

If you don't already have NFS:

helm repo add nfs-subdir-external-provisioner https://kubernetessigs.github.io/nfs-subdir-external-provisioner/

helm repo update

helm install nfs-server nfs-subdir-external-provisioner/nfs-subdir-external-provisioner \

```
--set nfs.server=<NFS SERVER IP> \
```

--set nfs.path=/exported/path \

--set storageClass.name=nfs-rwx

#### Step 2: Add Bitnami Jenkins Helm Repo

helm repo add bitnami https://charts.bitnami.com/bitnami helm repo update

#### Step 3: Prepare a jenkins-values.yaml File

controller:

replicaCount: 2

strategyType: RollingUpdate adminPassword:

"admin123"

persistence:

enabled: true





storageClass: "nfs-rwx"

accessMode: ReadWriteMany

size: 10Gi

#### healthProbes:

startupProbe:

enabled: true

initialDelaySeconds: 30

periodSeconds: 10

failureThreshold: 12

livenessProbe:

enabled: true

initialDelaySeconds: 90

periodSeconds: 10

failureThreshold: 6

#### service:

type: ClusterIP

#### ingress:

enabled: true

ingressClassName:

"nginx" hostname:

jenkins.local tls: false

#### O Step 4: Deploy Jenkins with Helm





helm install jenkins bitnami/jenkins -f jenkins-values.yaml -n jenkins --createnamespace

#### Step 5: Verify Jenkins Is Running in HA

kubectl get pods -n jenkins

You should see **two Jenkins pods** running. Check logs for successful startup:

kubectl logs <jenkins-pod-name> -n jenkins

O Step 6: Access Jenkins UI

Update /etc/hosts:

<INGRESS\_IP> jenkins.local Then

access: http://jenkins.local

Login with:

User: user Password:

admin123

- Step 7: Create a Sample Pipeline and Validate HA
  - 1. Create a simple Freestyle Job or Pipeline Job
  - 2. Trigger it
  - 3. While the job is running, **delete the active pod**:

kubectl delete pod <jenkins-0> -n jenkins

- ✓ Job should continue in the next pod
- UI remains accessible
- Pipeline state persists (thanks to shared NFS volume)
  - Step 8: Enable Horizontal Pod Autoscaling (Optional)





kubectl autoscale deployment jenkins -n jenkins --cpu-percent=60 --min=2 -- max=4

You can simulate CPU load to see HPA in action.

	What to Mention in Resume or Interview
=	Deployed Jenkins in High Availability mode using Kubernetes & NFS Ensured zero downtime CI/CD using multiple controller replicas and
	X PVC Configured liveness/startup probes for fault detection and HPA for scaling Simulated pod failure to validate pipeline continuity and HA setup

# **Project 3: Production-Grade HA Microservices Stack with Istio, Vault, MongoDB & ArgoCD on EKS**

Level: Advanced

**Solution** Objective:

Design and deploy a **highly available microservices platform** on **AWS EKS**, implementing:





#### HA MongoDB StatefulSets with data replication

- HA Vault setup with Integrated Raft storage
- Progressive Delivery via ArgoCD (with Canary rollouts)
- Traffic management and observability via Istio

#### **What You'll Learn**

- Build a real-world microservices architecture on EKS
- Configure HA for databases (MongoDB), secrets (Vault), and workloads
- Automate GitOps deployments using ArgoCD
- Enable secure service-to-service communication via Istio mTLS
- Validate high availability via failover simulations and health checks

#### High-Level Architecture

#### Prerequisites

- EKS Cluster (3 or more nodes across multiple AZs)
- Helm & kubectl installed
- Domain name & Route53 (for Ingress DNS)
- SSL certificate via cert-manager
- ArgoCD, Istio, and Vault Helm charts



#### O Step 1: Deploy Istio for Ingress & mTLS

istioctl install --set profile=demo -y
kubectl label namespace default istio-injection=enabled Deploy
Istio Gateway and VirtualService to expose your app: #
ingress-gateway.yaml
apiVersion: networking.istio.io/v1beta1 kind:
Gateway

Enable mTLS and DestinationRules to handle traffic securely.

#### Step 2: Deploy MongoDB as HA StatefulSet

To ensure **High Availability**, we will deploy MongoDB as a **StatefulSet** with:

- 3 replicas
- Pod-specific persistent volumes using EBS
- Headless service for internal DNS
- Readiness and liveness probes
- ReplicaSet configuration for automatic failover
- File 1: mongo-headless-service.yaml
  apiVersion: v1
  kind: Service
  metadata:

name: mongo

namespace: app

spec:



ports: - port: 27017 clusterIP: None # headless selector: app: mongo File 2: mongo-statefulset.yaml apiVersion: apps/v1 kind: StatefulSet metadata: name: mongo namespace: app spec: serviceName: "mongo" replicas: 3 selector: matchLabels: app: mongo template: metadata: labels: app: mongo spec: containers: - name: mongo

image: mongo:4.4





```
ports:
     - containerPort: 27017
    volumeMounts:
     - name: mongo-persistent-storage
      mountPath: /data/db
    livenessProbe:
     exec:
      command:
       - mongo
       - --eval
       - db.adminCommand('ping')
     initialDelaySeconds: 30
     periodSeconds: 10
    readinessProbe:
     exec:
      command:
       - mongo
       - --eval
       - db.adminCommand('ping')
     initialDelaySeconds: 10
     periodSeconds: 10
volumeClaimTemplates:
 - metadata:
   name: mongo-persistent-storage
  spec:
   accessModes: ["ReadWriteOnce"]
```



```
storageClassName: ebs-sc
    resources:
     requests:
      storage: 5Gi
File 3: mongo-init-config.yaml (for ReplicaSet Setup)
You need to initiate the ReplicaSet after the pods are running:
kubectl exec -it mongo-0 -n app -- mongosh
Then run:
js
rs.initiate({
 _id: "rs0",
 members: [
  { id: 0, host: "mongo-0.mongo.app.svc.cluster.local:27017" },
  { _id: 1, host: "mongo-1.mongo.app.svc.cluster.local:27017" },
  { _id: 2, host: "mongo-2.mongo.app.svc.cluster.local:27017" }
 ]
})
✓ Verify ReplicaSet Health
kubectl exec -it mongo-0 -n app -- mongosh
rs.status()
Look for:
      One PRIMARY
```

Two **SECONDARY** 

19



health: 1 for all members

#### **G** HA Testing

1. Kill the PRIMARY pod:

kubectl delete pod mongo-0 -n app

- 2. Run rs.status() again from mongo-1 or mongo-
- 2. Also create:
  - mongo-headless service (ClusterIP: None)
  - Readiness & Liveness probes
  - Affinity rules to spread pods across

AZs Use:

helm install mongodb percona/psmdb-db --namespace app --values values.yaml

#### O Step 3: Deploy Vault in HA Mode (Raft)

helm repo add hashicorp https://helm.releases.hashicorp.com helm repo update

helm install vault hashicorp/vault -n vault --create-namespace -f vault-values.yaml

Vault config (vault-values.yaml):

server:

ha:

enabled: true raft:

enabled: true

dataStorage:



- setWeight: 20



enabled: true storageClass: ebs-sc size: 10Gi injector: enabled: true Initialize and unseal Vault: kubectl exec -it vault-0 -n vault -- vault operator init Step 4: Deploy Microservice (e.g., NoteApp) # Deployment uses annotations for Vault injector annotations: vault.hashicorp.com/agent-inject: "true" vault.hashicorp.com/role: "noteapp" vault.hashicorp.com/agent-inject-secret-MongoDB\_\_\_ConnectionString: "secret/data/noteapp" ServiceAccount, Vault policy, and Kubernetes auth setup required. Deploy app with: kubectl apply -f noteapp-deployment.yaml O Step 5: Setup ArgoCD for GitOps + Canary kubectl create namespace argocd helm install argocd argo/argo-cd -n argocd Create Application YAML pointing to GitHub repo. Set up **Progressive Delivery**: canary: steps:



- pause: {}

- setWeight: 50

- pause: { duration: 30s }

#### **Step 6: Configure Ingress + TLS**

**⊘** Objective:

Secure external access to your NoteApp with:

- Istio Ingress Gateway
- DNS domain mapped to LoadBalancer
- TLS certificates managed by cert-manager

#### **Prerequisites:**

- Istio installed and configured
- cert-manager installed
   Install it via Helm if not already:

helm repo add jetstack https://charts.jetstack.io

helm repo update

helm install cert-manager jetstack/cert-manager \

--namespace cert-manager --create-namespace \

--set installCRDs=true

- A domain name (e.g., noteapp.example.com)
- A valid email address (used by Let's Encrypt)

☐ Step A: Create a ClusterIssuer (Let's Encrypt)

# cluster-issuer.yaml

apiVersion:

cert-manager.io/v1 kind:

#### ClusterIssuer





metadata: name: letsencrypt-prod spec: acme: email: your-email@example.com server: https://acme-v02.api.letsencrypt.org/directory privateKeySecretRef: name: letsencrypt-prod solvers: - http01: ingress: class: istio Apply it: kubectl apply -f cluster-issuer.yaml **Step B: Create Certificate Resource** # certificate.yaml apiVersion: cert-manager.io/v1 kind: Certificate metadata: name: noteapp-cert namespace: app spec: secretName: noteapp-tls issuerRef: name: letsencrypt-prod





kind: ClusterIssuer

commonName: noteapp.example.com

dnsNames:

- noteapp.example.com

Apply it:

kubectl apply -f certificate.yaml

This will automatically generate a TLS certificate and store it in the noteapp-tls Kubernetes secret.

☐ Step C: Configure Istio Gateway & VirtualService

# istio-gateway.yaml

apiVersion: networking.istio.io/v1beta1

kind: Gateway

metadata:

name: noteapp-gateway

namespace: app

spec:

selector:

istio: ingressgateway

servers:

- port:

number: 443

name: https

protocol: HTTPS

tls:

mode: SIMPLE

credentialName: noteapp-tls



Apply them:



# hosts: - noteapp.example.com # istio-virtualservice.yaml apiVersion: networking.istio.io/v1beta1 kind: VirtualService metadata: name: noteapp-vs namespace: app spec: hosts: - noteapp.example.com gateways: - noteapp-gateway http: - match: - uri: prefix: / route: - destination: host: noteapp port: number: 80





kubectl apply -f istio-gateway.yaml kubectl apply -f istio-virtualservice.yaml

#### Step D: DNS Configuration

1. Get LoadBalancer IP of Istio Gateway:

kubectl get svc istio-ingressgateway -n istiosystem

2. Add an A record in your DNS provider (Route53, GoDaddy,

etc.): noteapp.example.com → <LoadBalancer IP>

Wait for DNS propagation (usually <5 min).

**✓** Test Secure Access

Visit:

https://noteapp.example.com

- ✓ Should load the NoteApp securely
- ✓ TLS padlock visible
- ✓ Certificate issued by Let's Encrypt

What to Mention in Resume or Interview
<b>☑</b> Designed and implemented a production-grade HA microservices
platform on AWS EKS
☑ Enabled HA for Vault (Raft), MongoDB (ReplicaSet), and ArgoCD (GitOps
<b>✓</b> Secured service-to-service communication via Istio mTLS
✓ Used StatefulSets, PVs, and Helm charts to ensure self-healing and fault
tolerance
Applied Progressive Delivery strategy to minimize release risks

✓ Summary of Tools Used





Component	Tool
App Platform	Kubernetes on EKS
HA Ingress & Traffic	Istio
GitOps	ArgoCD
Secrets Management	Vault (HA Raft)
DB with Replication	MongoDB (StatefulSet + PVC)
Storage	AWS EBS
TLS Certs	cert-manager
CI/CD	Jenkins (Optional extension)

#### O Conclusion: Building Resilience

High Availability (HA) isn't just a technical checkbox — it's a **philosophy** that defines how you think, design, and operate systems at scale.

In this guide, you've walked through **three progressively challenging projects** that demonstrate your ability to:

- Architect resilient systems that **don't break under pressure**
- Handle failures gracefully using active-passive setups, replication, and self-healing deployments
- Integrate the latest in **GitOps, Service Mesh, StatefulSets, Vault HA**, and more
- Think like an **SRE**, **DevOps**, or **Platform Engineer** who owns not just uptime, but user trust





Whether you're prepping for your first job, switching roles, or trying to stand out in a crowded job market, showing that you've actually **implemented HA in real projects** is a **career-defining advantage**.

ιf	What	to Do	Next
----	------	-------	------

- Add these projects to your GitHub portfolio
- Mention HA highlights clearly in your resume
- De ready to explain these architectures in interviews
- Record short video demos or blog posts showcasing failover scenarios
- Start treating *availability* as a core design principle in all your future projects